

Functorial Composition over Task Decomposition: A Categorical Kernel for AI-Human Computation

Anonymous Author

mo:os Project

author@example.com

Multi-agent AI frameworks decompose tasks top-down, producing orchestrations that lack formal structural guarantees. We present `mo:os`, a runtime kernel realized as a symmetric monoidal category where computational entities are objects and state transitions are exactly four natural transformations (`ADD`, `LINK`, `MUTATE`, `UNLINK`) over a typed graph. Applying Lawvere’s functorial semantics, `mo:os` separates syntax (a JSON-schema-validated ontology of 21 object kinds and 16 morphism types) from semantics (a deterministic Go execution engine implementing a catamorphism $state(t) = fold(log[0..t])$). By restricting LLMs to emit typed morphisms rather than orchestration plans, complex behavior emerges from bottom-up operadic composition rather than unstructured decomposition. We formalize the Container Category, prove closure under parallel and sequential wiring, present the Recursive Semantic Bridge as a structure-preserving functor, construct a natural transformation from the kernel to the Model Context Protocol, and report implementation metrics from a working kernel: approximately 4,000 lines of Go, 118 nodes, 131 wires, 17 HTTP endpoints, an MCP bridge with 5 tool schemas, and deterministic replay from an append-only morphism log.

1 Introduction

The rapid adoption of large language models (LLMs) as orchestrators of multi-agent systems has exposed a fundamental architectural tension: the frameworks that decompose tasks—LangChain, CrewAI, AutoGen, and their successors—rely on the LLM’s embedding space to simultaneously serve as computation engine, state router, and correctness oracle. This conflation produces what we term the *semantic bottleneck*: as task complexity grows, recombination of sub-results depends on ad-hoc aggregation that is brittle to hallucination and context-window overflow.

We propose an inversion of control. Rather than asking the LLM to *orchestrate*, we restrict it to the role of a *Fuzzy Processing Unit* (FPU) that emits candidate morphisms over a strictly typed categorical state machine. The deterministic kernel accepts or rejects these morphisms against a semantic registry before committing them to an append-only log. This separation—syntax from semantics, in the sense of Lawvere [9]—yields structural guarantees that top-down decomposition cannot provide.

`mo:os` is an open-source (MIT) runtime kernel, implemented in Go, that realizes this architecture. Its categorical foundation draws on Fong and Spivak’s compositional framework [6], Spivak’s operadic wiring diagrams [15, 13], and Yau’s finite presentation results for wiring diagram operads [16]. Unlike Catlab (scientific computing) or Statebox (blockchain consensus) [4], `mo:os` targets the operational gap between LLM inference and industrial graph state management.

Contributions.

1. We define the *Container Category* $\mathcal{C}$ and prove it is closed under parallel and sequential wiring (Section 3).

2. We prove the *Authority Filtration* theorem and establish a Galois connection between the causal birth order and the stratum classification, showing that governance is a structural property of the morphism log (Section 3).
3. We formalize four invariant natural transformations and the operadic validation algebra that constrains them (Section 4).
4. We construct the *Recursive Semantic Bridge* as a structure-preserving functor $F : \mathbf{Syn} \rightarrow \mathbf{Sem}$ and show that hydration stages correspond to strata of realization (Section 5).
5. We formalize the LLM as an endofunctor rather than an orchestrator, and define a benchmark functor mapping providers to a metric space (Section 5).
6. We present *System 3* DAG reasoning as an extension of LogicGraph [1] (Section 6).
7. We construct the MCP bridge as a natural transformation from kernel categories to tool-schema categories and define pipeline complexity metrics with formal cost bounds (Section 7).
8. We report metrics from a working implementation: approximately 4,000 lines of Go, 118 nodes and 131 wires post-hydration, deterministic catamorphic replay, an MCP bridge with 5 tools, and a 21-kind ontology validated by an operad registry (Section 7).

2 Background: Functorial Semantics and Operadic Composition

2.1 Lawvere’s Functorial Semantics

Lawvere [9] established that an algebraic theory $\mathbf{T}$ can be separated from its models by defining models as product-preserving functors $M : \mathbf{T} \rightarrow \mathbf{Set}$. This separation of syntax (the theory) from semantics (the model) is the conceptual foundation of $\mathbf{mo} : \mathbf{os}$: the JSON ontology is our syntax category, and the Go kernel is a model that evaluates it.

2.2 Operadic Wiring Diagrams

Spivak [15] formalized wiring diagrams as an operad $\mathscr{W}$ where operations describe how boxes with typed ports are interconnected. Rupel and Spivak [13] extended this to temporal wiring diagrams with directed, length-bearing wires suitable for modeling dynamic information flow. Yau [16] proved that $\mathscr{W}$ admits a finite presentation with 8 generators and 28 relations.

In $\mathbf{mo} : \mathbf{os}$, the 16 morphism types of our ontology correspond to generators of a sub-operad of $\mathscr{W}$, and the semantic registry enforces the operadic relations: each wire must connect admissible port types with correct directionality.

2.3 Dynamic Operads and Agent Composition

Shapiro et al. [14] introduced *dynamic operads* via the monoidal double category $\mathbf{Org}$, showing how organized systems adapt their operadic structure in response to internal and external pressures. This framework supports the view of $\mathbf{mo} : \mathbf{os}$ as a dynamic categorical system where the graph topology evolves through morphism application while preserving algebraic invariants.

Bonchi, Pavlovic, and Sobocinski [5] extended functorial semantics from algebraic to relational theories, providing tools for reasoning about agent interactions as structure-preserving maps between relational networks—a pattern we exploit in the $\mathbf{mo} : \mathbf{os}$ wiring algebra.

3 The Container Category

Definition 1 (Container Category). *The Container Category $\mathcal{C}$ has:*

- **Objects:** typed nodes $n = (\text{urn}, \tau, \sigma, \pi)$ where urn is a unique identifier, $\tau \in \mathcal{T}$ is a type drawn from the ontology ($|\mathcal{T}| = 21$), $\sigma \in \{S_0, \dots, S_4\}$ is a stratum, and π is a payload.
- **Morphisms:** directed wires $w: n_1 \rightarrow n_2$ specified by the 4-tuple $(\text{src}, p_s, \text{tgt}, p_t)$ where p_s, p_t are typed ports.
- **Identity:** for each node n , the identity wire id_n is the trivial self-loop (elided in practice).
- **Composition:** wires compose when the target port of one matches the source port of the next, subject to the operadic constraints of the semantic registry.

The graph state at time t is the carrier of the catamorphism:

$$\text{state}(t) = \text{fold}(\log[0..t], \emptyset, \text{eval}) \quad (1)$$

where $\emptyset$ is the empty graph (initial object) and eval applies a single morphism envelope to produce a new state.

Definition 2 (Stratum Lattice). *The five strata form a partial order:*

$$S_0 \text{ (Authored)} \leq S_1 \text{ (Validated)} \leq S_2 \text{ (Materialized)} \leq S_3 \text{ (Evaluated)} \leq S_4 \text{ (Projected)}$$

Higher strata may depend on lower strata; lower strata must not depend on higher strata for meaning. S_4 projections are never ground truth.

Proposition 1 (Closure under Wiring). *Let G_1, G_2 be valid graph states in $\mathcal{C}$. The parallel composition $G_1 \otimes G_2$ (disjoint union of nodes and wires) and the sequential composition $G_1 \circ G_2$ (wire matching on shared ports) both yield valid graph states in $\mathcal{C}$, provided all wires satisfy the operadic constraints of the semantic registry.*

Proof sketch. Parallel composition preserves validity because disjoint union introduces no new port conflicts. Sequential composition requires that matched ports have compatible types and directions, which is enforced by the registry’s `ValidateLink` function checking source port direction (“out”), target port direction (“in”), and target type membership in the source port’s admissible target set. $\square$

3.1 Ownership as Full Subcategory

The OWNS morphism (*MOROI*) induces a forest structure on $\mathcal{C}$. For any actor a , the *scope* of a is the full subcategory $\mathcal{C}_a \hookrightarrow \mathcal{C}$ consisting of all nodes transitively reachable from a via OWNS wires. This is computed by breadth-first search under a read lock, yielding all internal wires—those whose source and target both lie in $\mathcal{C}_a$.

3.2 Authority Filtration

The morphism `log` and the stratum lattice each induce a filtration of $\mathcal{C}$, but these two filtrations are *independent* and run in opposite directions on the lower strata. Making this precise is the key structural result of the paper.

Definition 3 (Causal Filtration). *Let the morphism log have t entries. For each $k \in \{0, \dots, t\}$, let $\mathcal{C}_{\leq k}$ be the full subcategory of $\mathcal{C}$ whose objects have been created by ADD envelopes at log positions $\leq k$. The chain*

$$\emptyset = \mathcal{C}_{\leq -1} \hookrightarrow \mathcal{C}_{\leq 0} \hookrightarrow \mathcal{C}_{\leq 1} \hookrightarrow \dots \hookrightarrow \mathcal{C}_{\leq t} = \mathcal{C}$$

is the causal filtration of $\mathcal{C}$. Each step adds at most one object (the node created by the ADD at position k , if any).

Definition 4 (Stratum Filtration). *For each stratum S_i , let $\mathcal{C}_{\leq i}^\sigma$ be the full subcategory of $\mathcal{C}$ whose objects have stratum $\leq S_i$. The chain*

$$\mathcal{C}_{\leq 0}^\sigma \hookrightarrow \mathcal{C}_{\leq 1}^\sigma \hookrightarrow \mathcal{C}_{\leq 2}^\sigma \hookrightarrow \mathcal{C}_{\leq 3}^\sigma \hookrightarrow \mathcal{C}_{\leq 4}^\sigma = \mathcal{C}$$

is the stratum filtration of $\mathcal{C}$ by full subcategories.

The kernel self-seeds at boot: the first morphism in every log is $\text{ADD}(\text{urn}:\text{moos}:\text{kernel}:\text{wave}-0)$, a node at stratum S_2 . Every subsequent node enters $\mathcal{C}$ via a finite sequence of ADD and LINK morphisms traceable back to this seed through the *causal* filtration—not the stratum filtration.

Theorem 1 (Authority Filtration). *Every node $n \in \text{Ob}(\mathcal{C})$ has a finite causal path in the morphism log to the kernel seed. Formally: for each n with $\text{birth}(n) = k > 0$, there exists a LINK morphism $\ell: m \rightarrow n$ in the log with $\text{birth}(m) < k$ —that is, the authorizing node m already exists in $\mathcal{C}_{\leq k-1}$. By induction, every node is connected to the seed by a finite chain of such LINK morphisms.*

Proof sketch. By construction of the boot sequence and the hydration pipeline. The kernel seed at S_2 is the first ADD (log position 0). Source nodes at S_1 are seeded next via $\text{ADD} + \text{LINK}(\text{kernel}, \text{source})$. Industry entities at S_0 enter later via $\text{ADD} + \text{LINK}(\text{source}, \text{industry})$. Instance nodes at S_2 are hydrated via $\text{ADD} + \text{LINK}(\text{industry}, \text{instance})$ or $\text{LINK}(\text{kernel}, \text{instance})$. Projections at S_4 are functor outputs. At each step, the LINK envelope’s precondition (both endpoints must exist) guarantees that the authorizing node has a strictly earlier birth position. Since the log is finite, the inductive chain terminates at the seed. $\square$

Crucially, the authorizing LINK need not respect the stratum order: the seed at S_2 authorizes source nodes at S_1 and industry entities at S_0 . Authority flows *down* the stratum lattice while semantic dependency flows *up*. This contravariance is not a defect but a structural feature, which we now make precise.

Proposition 2 (Contravariant Orders). *Let complete: $\Sigma \rightarrow \mathbb{N}$ map each stratum S_i to the last log position at which a node of stratum S_i is created, and let filled: $\mathbb{N} \rightarrow \Sigma$ map each log position k to the highest stratum S_i such that all S_i -nodes have birth $\leq k$. Then filled $\dashv$ complete is a Galois connection between $(\mathbb{N}, \leq)$ and $(\Sigma, \leq)^{\text{op}}$:*

$$\text{filled}(k) \geq S_i \iff k \geq \text{complete}(S_i).$$

In particular, the stratum order and the causal birth order are contravariant on the lower strata: the kernel seed is causally initial (birth = 0) but resides at S_2 , while S_0 industry entities are causally posterior but immutable.

Proof sketch. Both maps are well-defined on a finite log. The equivalence holds by definition: “stratum S_i is fully present by depth k ” iff “ k is at least as deep as the last S_i creation.” Contravariance on the lower strata follows from the boot order: $\text{complete}(S_2) = 0 < \text{complete}(S_1) < \text{complete}(S_0)$, since the S_2 seed precedes S_1 sources, which precede S_0 industry entities in the log. The `applyMutate` guard (S_0 nodes reject `MUTATE`) enforces the stratum constraint independently of the causal order: a node that is causally young can still be immutable. $\square$




Figure 1: Contravariant orders. *Left*: the causal filtration ordered by log position (Def. 3). *Right*: the stratum filtration ordered by type classification (Def. 4). Dashed cross-links map each causal stage to its stratum, revealing the contravariance: the kernel seed is causally first but resides at S_2 , while S_0 industry entities are causally posterior but immutable. The pair (filled, complete) forms a Galois connection (Prop. 2) between these two independent orders.

Together, Theorem 1 and Proposition 2 upgrade Axiom 5 (“governance is structural”) from a design principle to a pair of formal results. Authority is not enforced by middleware or access control lists, but by the causal structure of the morphism log itself. The Galois connection provides a formal tool for boot sequencing: $\text{complete}(S_i)$ tells the kernel when stratum S_i is fully loaded, and $\text{filled}(k)$ tells it what structural constraints are satisfiable at any point during replay. The contravariance between the two orders is analogous to Wolfram’s distinction between the causal graph (on rewriting events) and the spatial graph (on elements): authority lives on events, strata live on elements (Figure 1).

4 The Semantic Bottleneck and Invariant Morphisms

4.1 The Semantic Bottleneck

The prevailing paradigm for multi-agent orchestration relies on the LLM’s embedding space as both computation engine and state router. We formalize this structural failure as the *semantic bottleneck*: a system architecture where semantic inference and operational state management share a single, untyped channel. As task complexity grows, this channel saturates, producing hallucinated execution paths and brittle error recovery.




Figure 2: The semantic bottleneck (left): LLMs serve as both orchestrator and aggregator through an untyped channel, producing brittle compositions. Functorial composition (right): the LLM acts as a Fuzzy Processing Unit emitting typed morphism envelopes, validated by the operad registry before commitment to the append-only log.

4.2 Four Invariant Natural Transformations

`mo:os` resolves the bottleneck by restricting all state mutations to exactly four natural transformations:

NT	Signature	Semantic	Guard
ADD	$\emptyset \rightarrow \text{Node}$	Create node	Duplicate URN $\Rightarrow \text{ErrNodeExists}$
LINK	$N \times N \rightarrow \text{Wire}$	Create wire	Operadic port validation
MUTATE	$N \rightarrow N$	Update payload	Version CAS + S_0 block
UNLINK	$\text{Wire} \rightarrow \emptyset$	Remove wire	Wire not found $\Rightarrow \text{error}$

These four operations are *complete*: any graph state reachable from the empty graph can be expressed as a finite composition of ADD, LINK, MUTATE, and UNLINK envelopes. The LLM (as FPU) proposes candidate envelopes; the kernel validates and commits them.

4.3 Operadic Validation

The semantic registry $\mathcal{R}$ is an operad algebra over the port structure of $\mathcal{C}$. For each type $\tau \in \mathcal{T}$, the registry specifies:

- **Allowed strata**: the set of valid strata for instances of τ .
- **Port specifications**: for each named port, its direction (in/out) and admissible target types.
- **Mutability**: whether instances admit MUTATE.

A LINK envelope $w: n_1 \xrightarrow{(p_s, p_t)} n_2$ is valid if and only if: (1) p_s exists on $\tau(n_1)$ with direction “out”, (2) p_t exists on $\tau(n_2)$ with direction “in”, and (3) $(\tau(n_2), p_t)$ is in the admissible target set of p_s . This three-part check implements the operadic composition constraints of $\mathcal{W}$.




Figure 3: The hydration pipeline as stratum promotion. Each stage corresponds to a kernel operation: schema validation ($S_0 \rightarrow S_1$), materialization into Programs ($S_1 \rightarrow S_2$), fold evaluation ($S_2 \rightarrow S_3$), and scoped projection via BFS ($S_3 \rightarrow S_4$). S_4 outputs are never ground truth.

5 The Recursive Semantic Bridge

We formalize the `mo:os` architecture as a realization of Lawvere’s functorial semantics. The database graph serves as the abstract syntax tree (AST) of our language.

Definition 5 (Syntax Category). **Syn** has objects = Container types (the 21 kinds of the ontology) and morphisms = wiring diagrams (the 16 morphism types).

Definition 6 (Semantics Category). **Sem** has objects = Go runtime states and morphisms = deterministic kernel evaluations (the fold step of Equation 1).

The *Recursive Semantic Bridge* is the structure-preserving functor:

$$F: \mathbf{Syn} \rightarrow \mathbf{Sem} \quad (2)$$

$$\begin{array}{ccc}
 \mathbf{Syn} & \xrightarrow{F} & \mathbf{Sem} \\
 \alpha \downarrow & & \downarrow F(\alpha) \\
 \mathbf{Syn}' & \xrightarrow{F} & \mathbf{Sem}'
 \end{array} \quad (3)$$

where α represents a morphism (state transition) in the syntax, and $F(\alpha)$ is the corresponding Go kernel evaluation. The functor preserves composition: applying two morphisms sequentially in **Syn** yields the same result as applying their F -images sequentially in **Sem**.

5.1 Hydration as Stratum Promotion

The hydration pipeline maps directly onto the stratum lattice:

Stage	Stratum	Kernel Operation
Authored	S_0	JSON instance files on disk
Validated	S_1	Schema check (JSON Schema draft-07)
Materialized	S_2	Materialize: instance $\rightarrow$ Program
Evaluated	S_3	ApplyProgram: fold into graph state
Projected	S_4	ScopedSubgraph: BFS lens

The key invariant is that S_4 projections are *never* ground truth: they are views derived from the evaluated state, and the graph can always be reconstructed from the morphism log via replay (Equation 1).




Figure 4: The three-layer tower. Industry data enters via the translation functor $\mathcal{F} : \mathbf{Syn}_{BN} \rightarrow \mathbf{Syn}_{LP}$. The superset ontology serves as the free category whose models (FUN01–FUN05) are structure-preserving functors to target categories. Stable patterns are promoted into the kernel via ι_K ; the kernel exposes its structure back to the graph via ρ_K (the adjoint restriction). File paths ground each layer in the codebase.

5.2 Harness Agnosticism

Because the functor F maps syntax to semantics uniformly, any external tool—an LLM, a database, a UI framework—is merely an object in **Syn** whose interaction with the graph is governed by F . This achieves *harness agnosticism*: the UI surface is a direct, consistent projection of the categorical state, requiring no intermediary state-management synchronization. Nested sub-workspaces are submodules in the categorical composition, inheriting structural integrity via algebraic closure.

5.3 LLM as Endofunctor

In `mo:os`, an LLM is not an orchestrator but an object of type `agnostic_model` in $\mathcal{C}$. Its interface to the rest of the graph is mediated entirely by `CAN_ROUTE` morphisms connecting it to `protocol_adapter` nodes. This embedding has a precise categorical reading: the LLM acts as an endofunctor on the subcategory of morphism proposals, accepting candidate envelopes and returning refined candidates—all subject to the kernel’s validation.

The *benchmark functor* F_{bench} formalizes model selection as a functorial problem. Let $\mathcal{C}_{Prov}$ be the full subcategory of $\mathcal{C}$ whose objects are of type `provider` and whose morphisms are feature-bearing wires (e.g., `CAN_ROUTE`, `BENCHMARKS_ON`).

Definition 7 (Benchmark Functor). *The benchmark functor $F_{bench} : \mathcal{C}_{Prov} \rightarrow \mathbf{Met}$ maps:*

- Each provider object p to a point in the metric space $(\mathbb{R}^k, d)$ representing its benchmark score vector (latency, accuracy, cost, throughput).
- Each feature morphism $f : p_1 \rightarrow p_2$ to a non-expansive map preserving the partial order on scores.

Model selection then reduces to choosing the provider whose F_{bench} -image optimizes a given objective. This is the *collider pattern*: models are benchmarked agnostically as objects in the graph, their capabilities are morphisms, and composition of model pipelines happens categorically through the typed wiring algebra—not through prompt engineering or ad-hoc orchestration scripts.




Figure 5: Layered architecture of the `mo:os` kernel. The pure core (`cat` + `fold`) contains zero IO imports. The effect shell wraps it with `sync.RWMutex` for concurrent access. Both LLM agents and human actors interact exclusively through typed morphism envelopes.

6 System 3: Multi-Path DAG Reasoning

We extend `LogicGraph` [1] to define System 3 reasoning over the morphism space:

- **System 1:** Intuitive token prediction (standard LLM inference).
- **System 2:** Linear chain-of-thought (sequential reasoning).
- **System 3:** DAG exploration over typed morphism space.

In System 3, the deterministic kernel spawns candidate sub-graphs where each edge is a structure-preserving morphism validated against the `operad` registry. By applying tree search heuristics (MCTS or beam search) over the morphism space, `mo:os` simulates branching future states. Because each branch is restricted to valid categorical syntax, the search space is automatically pruned of structurally inconsistent outcomes.

This enables high-stakes reasoning—infrastructure deployment, data migration, multi-step planning—where the error surface of a single LLM prompt would be prohibitively large. The typed morphism space provides a natural pruning function: invalid compositions are rejected before evaluation, and the catamorphic replay guarantees that any accepted path can be audited.

7 Implementation

`mo:os` is implemented in Go (1.22+) as a single-binary kernel with zero external dependencies beyond the standard library. The architecture separates pure computation from effectful state management:

Layer	Package	Responsibility
Types	cat	Core types: Node, Wire, Graph-State, Envelope, Program. 4 morphism types, 5 strata.
Pure fold	fold	Catamorphism: Evaluate, Replay. Zero IO imports.
Operad	operad	Semantic registry: 21 Type-Specs with port/direction/target validation.
Effect shell	shell	Runtime: RWMutex-guarded state, Apply, SeedIfAbsent.
Transport	transport	17 HTTP routes on port 8000. Embedded Explorer UI via go:embed (transport/static/).
Hydration	hydration	Materialize: instance JSON $\rightarrow$ Program. HydrateAll: batch boot.
Lens	lens	Composable graph filter: Apply(spec, state), 5 predicate types (kind, stratum, category, port, neighborhood), scope composition.
MCP bridge	mcp	MCP server on port 8080: SSE transport, JSON-RPC 2.0, 5 tool schemas.
Functor	functor	Benchmark functor FUN05: $\mathcal{C}_{\text{Prov}} \rightarrow \mathbf{Met}$.

7.1 Catamorphic Replay

The morphism log is an append-only JSONL file. State reconstruction is a left fold:

Listing 1: Pure replay (simplified)

```
func Replay(log []Envelope) GraphState {
    state := NewGraphState() // empty
    for _, env := range log {
        state = Evaluate(state, env)
    }
    return state
}
```

The fold package imports no IO packages (os, net, sync); it operates purely on values. The effect shell (shell.Runtime) wraps the pure fold with sync.RWMutex for concurrent read access and exclusive write access.

7.2 Idempotent Seeding

Boot uses `SeedIfAbsent`, which absorbs `ErrNodeExists` and `ErrWireExists`. This makes repeated boots idempotent: the first boot creates 13 seed morphisms; subsequent boots replay the log and produce zero new entries.

7.3 MCP Bridge as Natural Transformation

The Model Context Protocol (MCP) bridge [2] exposes the kernel to LLM tool-use via SSE on port 8080. We formalize this bridge as a natural transformation between two categories.

Definition 8 (Kernel and Tool Categories). *Let $\mathcal{C}_K$ be the category whose objects are graph operations (ADD, LINK, MUTATE, UNLINK, and the read operations STATE, LOOKUP, SCOPE) and whose morphisms are kernel evaluations (applications of the catamorphic fold). Let $\mathcal{C}_L$ be the category whose objects are JSON tool schemas and whose morphisms are JSON-RPC 2.0 invocations.*

Definition 9 (MCP Natural Transformation). *The MCP bridge is a natural transformation $\eta : \mathcal{C}_K \Rightarrow \mathcal{C}_L$ whose five components are:*

Component η_i	Kernel operation $\rightarrow$ Tool schema
η_{state}	STATE $\mapsto$ <i>graph_state</i>
η_{lookup}	LOOKUP $\mapsto$ <i>node_lookup</i>
η_{apply}	APPLY $\mapsto$ <i>apply_morphism</i>
η_{scope}	SCOPE $\mapsto$ <i>scoped_subgraph</i>
η_{bench}	BENCHMARK $\mapsto$ <i>benchmark_project</i>

The naturality condition states that for any kernel morphism $f : A \rightarrow B$ in $\mathcal{C}_K$:

$$\begin{array}{ccc}
 \mathcal{C}_K(A) & \xrightarrow{\eta_A} & \mathcal{C}_L(\eta_A) \\
 f \downarrow & & \downarrow \eta(f) \\
 \mathcal{C}_K(B) & \xrightarrow{\eta_B} & \mathcal{C}_L(\eta_B)
 \end{array} \tag{4}$$

Concretely: applying a morphism in the kernel and then projecting the result via an MCP tool yields the same output as first projecting via the tool and then applying the corresponding JSON-RPC invocation. This ensures that LLM clients observe a consistent view of the kernel state regardless of whether they query before or after a mutation is committed.

7.4 Pipeline Complexity and Subgraph Metrics

The scoped subgraph operation (Section 3) enables us to define formal complexity metrics for pipeline reasoning.

Definition 10 (Subgraph Cost Model). *For an actor a with scoped subgraph $G_a = \text{ScopedSubgraph}(a)$, define:*

- Discovery cost: $D(G_a) \propto |E(G_a)|$ — the cost of scanning edges to discover structural relationships.
- Retrieval cost: $R(G_a) \propto |V(G_a)|$ — the cost of accessing node state within the scope.

- Complexity ratio: $\kappa(G_a) = |E(G_a)|/|V(G_a)|$ — the average edge density per node, measuring wiring complexity.

The BFS scoping algorithm computes G_a in $O(|V| + |E|)$ time under RLock, where $|V|$ and $|E|$ are the sizes of the full graph. Once scoped, all subsequent operations on G_a operate on the reduced subgraph.

Theorem 2 (Discovery Cost Reduction). *Let E_{total} be the edge set of the full graph and $E_{\text{scope}} = E(G_a) \subseteq E_{\text{total}}$ be the edge set of the scoped subgraph induced by the ownership functor. The transition from task decomposition (which must inspect all edges to discover relevant structure) to functorial composition via scoped projection reduces discovery cost from $O(|E_{\text{total}}|)$ to $O(|E_{\text{scope}}|)$.*

Proof sketch. Task decomposition frameworks lack an ownership functor and must traverse the full graph to determine which edges are relevant to a given actor. The ownership functor $\Omega: \mathcal{C} \rightarrow \mathbf{Forest}$ partitions the graph into disjoint subtrees. For any actor a , the scoped subgraph $G_a = \Omega^{-1}(a)$ contains exactly those edges internal to a ’s ownership subtree. Since Ω is computed once during BFS and the result is cached, subsequent discovery within G_a examines only $|E_{\text{scope}}|$ edges. The ratio $|E_{\text{scope}}|/|E_{\text{total}}|$ is bounded by the ownership fraction of the actor, which in practice is $\ll 1$ for large graphs. $\square$

7.5 Data Lens: Functorial Projection as a User Interface

The `lens` package materializes functorial projection as a composable query interface. A *lens specification* $L = (\rho_1, \dots, \rho_n, m)$ pairs a sequence of predicates over kind, stratum, category, port, or k -hop neighbourhood with a composition mode $m \in \{\cap, \cup\}$. For scoped queries the kernel first applies the ownership functor Ω (Section 3), then applies Λ (lens filter): the composed functor $\Lambda \circ \Omega: \mathcal{C} \rightarrow \mathcal{C}_{\text{filtered}}$ is served via `GET /state/lens` and `POST /state/lens`.

The Explorer UI renders the filtered `GraphState` as structured tables grouped by object kind, morphism type, type registry, and morphism history. This design makes categorical structure legible through tabular projection rather than spatial layout, demonstrating that functorial composition provides auditable structure *without* requiring a dedicated hypergraph renderer.

7.6 Metrics

Metric	Value
Go source lines	~4,000
Ontology object kinds	21
Ontology morphism types	16
Invariant natural transformations	4
Strata	5
Categories	22
Functors	5
HTTP routes	17
MCP tools	5
Nodes (post-hydration)	118
Wires (post-hydration)	131
Test packages (all green)	9
External dependencies	0

8 Related Work

Categorical frameworks. Catlab [12] provides Julia libraries for computational category theory; Statebox [4] applies Petri nets to blockchain consensus; Lynch and Breiner [10] implement nested wiring diagrams in Semagrams. None target the operational gap between LLM inference and runtime graph state management.

Multi-agent orchestration. LangChain, CrewAI, AutoGen, and LangGraph rely on LLM-driven task decomposition. These frameworks lack formal structural guarantees: correctness depends on the LLM maintaining coherent state across unbounded context windows.

Operadic composition. Spivak’s wiring diagram operads [15, 13] and Yau’s finite presentation results [16] provide the mathematical foundation for our port-typed composition. Shapiro et al.’s dynamic operads [14] formalize how operadic structure evolves—a pattern mo:os realizes through its mutable graph topology.

Functorial semantics for computation. Bonchi et al. [5] extend functorial semantics to relational theories. Gu and Zanasi [8] apply functorial semantics as a unifying framework for logic programming. mo:os applies these ideas to AI-human computation, where the “programs” are morphism envelopes emitted by both human actors and LLM agents.

Category theory for AI. Pan [11] proposes a categorical framework at the token level of transformer architectures. Gavranović et al. [7] develop an algebraic theory of deep learning architectures using 2-categories of parametric maps. These works formalize AI *internals* categorically; mo:os instead formalizes the *interface* between AI agents and shared state.

Neuro-symbolic reasoning. LogicGraph [1] benchmarks multi-path reasoning in high-dimensional search spaces. mo:os ’s System 3 extends this by embedding DAG reasoning within a categorically typed morphism space, providing structural pruning guarantees absent in untyped search.

9 Conclusion

We have demonstrated that the transition from top-down task decomposition to bottom-up functorial composition provides a rigorous foundation for AI-human computation. By anchoring LLMs as coprocessors within a typed categorical kernel, mo:os resolves the semantic bottleneck while preserving the expressiveness needed for complex multi-agent workflows.

The working implementation—approximately 4,000 lines of Go with zero dependencies, 21 object kinds validated by an operad registry, 118 nodes and 131 wires post-hydration, and deterministic categorical replay—shows that mathematical rigor is not an obstacle to practical systems but rather their most robust enabling technology. The authority filtration (Theorem 1) and contravariant orders (Proposition 2) upgrade the governance axiom from a design principle to a pair of formal results: every node’s provenance is auditable by tracing LINK morphisms back through the causal filtration to the kernel seed, and the Galois connection between causal birth order and stratum classification provides a formal tool for boot sequencing and structural consistency verification. The MCP bridge (Section 7.3) realizes the

natural transformation $\eta: \mathcal{C}_K \Rightarrow \mathcal{C}_L$, making the kernel’s categorical operations directly accessible to any MCP-compliant LLM client via five tool schemas over SSE and JSON-RPC 2.0. The benchmark functor $F_{\text{bench}}: \mathcal{C}_{\text{Prov}} \rightarrow \mathbf{Met}$ (Definition 7) reduces model selection to a functorial optimization problem, and the pipeline complexity metrics (Theorem 2) formalize the cost advantage of scoped projection over full-graph traversal.

Future work includes: (1) upgrading the transport layer to HTTP/3 via quic-go [3], mapping stratum-specific reliability requirements to independent QUIC streams; (2) pipeline complexity metrics based on the subgraph $|E|/|V|$ ratio, enabling automatic detection of over-wired subsystems; (3) LLM-as-morphism formalization, embedding `agnostic_model` inference as a first-class endomorphism in the typed graph; and (4) quantitative comparison of morphism-based error recovery rates against traditional LLM orchestration frameworks such as LangChain and AutoGen.

References

- [1] Anonymous (2025): *Benchmarking Multi-Path Logical Reasoning in High-Dimensional Search Spaces*. Internal Report.
- [2] Anthropic (2024): *Model Context Protocol (MCP)*. <https://modelcontextprotocol.io>.
- [3] Mike Bishop (2022): *HTTP/3*. RFC 9114. Available at <https://datatracker.ietf.org/doc/html/rfc9114>.
- [4] Jelle Herold Boer et al. (2019): *Mathematical Specification of the Statebox Language*. In: *Proceedings of ACT 2019*.
- [5] Filippo Bonchi, Dusko Pavlovic & Pawel Sobocinski (2017): *Functorial Semantics for Relational Theories*. arXiv preprint arXiv:1711.08699.
- [6] Brendan Fong & David I Spivak (2018): *Seven Sketches in Compositionality: An Invitation to Applied Category Theory*. Cambridge University Press.
- [7] Bruno Gavranović, Paul Lessard, Andrew Duber & Daniel Murfet (2024): *Categorical Deep Learning: An Algebraic Theory of Architectures*. arXiv preprint arXiv:2402.15332.
- [8] Tao Gu & Fabio Zanasi (2021): *Functorial Semantics as a Unifying Perspective on Logic Programming*. In: *9th Conference on Algebra and Coalgebra in Computer Science (CALCO 2021)*, Schloss Dagstuhl – Leibniz-Zentrum für Informatik, pp. 17:1–17:22.
- [9] F. William Lawvere (1963): *Functorial Semantics of Algebraic Theories*. *Proceedings of the National Academy of Sciences* 50(5), pp. 869–872.
- [10] Owen Lynch & Spencer Breiner (2023): *Nested wiring diagrams in Semagrams*. In: *Proceedings of ACT 2023, EPTCS* 397.
- [11] Wuming Pan (2024): *Token Space: A Category Theory Framework for AI Computations*. arXiv preprint arXiv:2404.11624.
- [12] Evan Patterson, Andrew Baas & James Fairbanks (2024): *Catlab.jl: A Framework for Applied Category Theory in Julia*. <https://github.com/AlgebraicJulia/Catlab.jl>.
- [13] Dylan Rupel & David I Spivak (2013): *The operad of temporal wiring diagrams: formalizing a graphical language for discrete-time processes*. arXiv preprint arXiv:1307.6894.
- [14] Brandon T Shapiro & David I Spivak (2022): *Dynamic Operads, Dynamic Categories: From Deep Learning to Prediction Markets*. arXiv preprint arXiv:2205.03906.
- [15] David I Spivak (2013): *The operad of wiring diagrams: formalizing a graphical language for databases, plug-ins, and concurrent routing*. arXiv preprint arXiv:1305.0297.
- [16] Donald Yau (2018): *Operads of Wiring Diagrams*. *Lecture Notes in Mathematics* 2192, Springer. ArXiv:1512.01602.